

CSE 576 Project 3: Generative Discriminative Image Classifier.

Brendan Miller

January 13, 2016

1 Introduction

For this project I built a classifier that recognizes images that contain certain classes of objects. To do this I use the generative/discriminative framework discussed in Yi Li[1].

Whereas Li's paper used features extracted from clustered regions such as color and texture, this project also uses descriptors extracted from keypoints. Specifically, the keypoints and descriptors generated by SIFT.

2 Related work

This project is primarily based on *A Generative/Discriminative Learning Algorithm for Image Classification* by Yi Li[1]. Li's generative/discriminative classifier framework works in two stages.

First, features are extracted from a training set of positive examples. From this set of features, the expectation maximization algorithm is used to create a gaussian mixture model of the distribution of the features.

For an object o and a feature type a the probability of a particular feature vector X^a appearing can be calculated from the gaussian mixture model as

$$P(X^a|o) = \sum_{m=1}^{M^a} w_m^a N(X^a; \mu_m^a, \Sigma_m^a)$$

where M^a is the number of clusters in the mixture model for feature a . For the m th gaussian of feature a w_m^a is the weight, μ_m^a is the mean, and Σ_m^a is the covariance matrix.

In the discriminative step, both positive and negative examples are used to train a classifier, in this case a pretty standard 3 layer neural network.

For the i th image the joint probabilities of the r th feature in the image and the m th component gaussian of feature type a is calculated like so:

$$P(X_{i,r}^a, m^a) = w_m^a N(X_{i,r}^a, \mu_m^a, \Sigma_m^a)$$

to determine how well an image I_i matched a component gaussian m_a , I just use the maximum for all features in the image. My reasoning is that having some background features that do not match, should reduce the score of the overall match:

$$P(I_i, m^a) = \max(\{P(X_{i,r}^a, m^a) | r \in \text{features of type } a \text{ in the } i\text{th image}\})$$

For each image I_i all of these maximum joint probability values $P(I_i, m^a)$ for all components m^a and all features a are aggregated into a vector. This vector is used to train a 3 layer neural network. The neural network is trained to return 1 for vectors from positive example images, and 0 for negative examples.

In practice, given a vector of maximum joint probabilities, the neural network will return a value roughly between 0 and 1 that corresponds to the confidence that object o is present. For this reason, a threshold over the output of the neural network can be varied to trade off the likely number of false positives and false negatives returned. In my implementation I use a threshold of 0.5.

3 Your method

My implementation keeps the underlying framework of Li's paper, but in addition to using regional features mine also uses descriptors from key-points.

In my implementation for each image I extracted SIFT[2] descriptors using the OpenCV library and fed these into the generative discriminative framework discussed in the last section.

I found that using a fairly large number of component gaussians in the SIFT mixture model helped greatly with classification accuracy. I believe this is because a given object class can have many distinctive parts, and a component gaussian is needed to represent the distribution of features in each part.

For instance, in a mixture model representing the distribution of SIFT features extracted from motorcycle images, you would want one component gaussian to represent the descriptors extracted from headlamps, another to represent wheels, another for tailpipes, and so on.

Using only SIFT features and a gaussian mixture model with 50 components, I could achieve classification accuracy around 75% to 80%.

The number of gaussian components can continue to be increased to increase accuracy. I capped it at 50 for most of my testing mainly because the computational overhead in the training phase became significant with higher numbers of gaussians.

In addition to SIFT descriptors, I also ran kmeans on each image to extract color clusters in CIE L*a*b* color space, the mean color each of which I fed into the generative/discriminative framework.

Color is mainly useful for classifying objects that have a distinctive color, such as faces which contain the human skin color. They were less useful for classifying planes, and motorcycles, all of which can be painted a variety of colors. Cars were well classified by color; however as I discuss in the experimental results I think this has more to do with the uniformity of the car data set than anything else.

4 Experiments and results

I used motor bikes, faces, airplanes, and cars from the caltech image sets linked from the original project description.

Links to datasets used for testing:

```
http://www.vision.caltech.edu/Image_Datasets/cars_brad/cars_brad.tar
http://www.vision.caltech.edu/Image_Datasets/motorbikes_side/motorbikes_side.tar
http://www.vision.caltech.edu/Image_Datasets/airplanes_side/airplanes_side.tar
http://www.vision.caltech.edu/Image_Datasets/faces/faces.tar
```

I use a training set of 400 images and a test set of 200 images randomly selected from a larger collection. Each set contains 50/50 positive and negative examples.

For the data collected here I used a gaussian mixture model with 10 components to represent the color distribution of each class and a mixture model with 50 components to represent the SIFT feature distribution.

4.1 SIFT features only

Using SIFT alone performance was pretty decent across a range of image types. In general increasing the number of gaussians in the mixture model improves SIFT at the expense of additional training time.

Image Type	Accuracy	False Positives	False Negatives
Cars	0.735	0.125	0.14
Planes	0.775	0.09	0.135
Faces	0.78	0.13	0.09
Bikes	0.755	0.11	0.135

4.2 CIE L*a*b* color features only

Not surprisingly faces do a little bit better than planes and motorcycles when only using color for image classification, as faces have a fairly distinctive color tone that has been shown to be well modeled by a gaussian distribution. By comparison planes and motorcycles can be painted a variety of arbitrary colors.

What is most surprising is that cars are classified so well by color alone. My suspicion is that this is because the training images shows cars as they are driving surrounded by pavement, and that the color of the pavement is really what the classifier is finding. This belief is given further evidence because many false positives were non-car images that had pavement in them.

For this reason I consider the motorcycle data a lot better representation of the power of this model, as it has motorcycles in a variety of scenes, including scenes where the background has been cropped out.

Image Type	Accuracy	False Positives	False Negatives
Cars	0.905	0.085	0.01
Planes	0.77	0.075	0.155
Faces	0.795	0.095	0.11
Bikes	0.725	0.19	0.085

4.3 Combined SIFT and CIE L*a*b* color features

This set of results really shows the strength of the discriminative generative model. Every result is a bit better than the results from the individual features. Face classification in particular really benefits from a combination of both color and keypoint descriptor data.

Image Type	Accuracy	False Positives	False Negatives
Cars	0.94	0.035	0.025
Planes	0.78	0.11	0.11
Faces	0.9	0.055	0.045
Bikes	0.805	0.08	0.115

This is an example of a correctly classified plane.



False positives: It's a little tricky to interpret why this example was classified as a plane, but my guess is that the white color of the car, which is more common on commercial aircraft than in automobiles, partially had to do with it.



False negatives: Here's an example of a false negative during face classification. This image lacks the distinctive color of skin, and is also hand drawn so keypoint descriptors may not be quite the same as with real faces.



5 Future work

One idea to extend this project would be to try to work backwards from an image classified as having an object, to try to find which features were responsible for the positive classification.

If some book-keeping were done to keep track of the locations and size of the original features, a bounding box around the region likely containing the object could be found.

One obstacle to this approach is that the fairly complex structure of a 3 layer neural network is a little tricky to interpret when trying to determine which input data is responsible for the result. For that reason experiments with either a 2 layer network, or some other simpler classifier could be done.

My intuition is that a 3 layer network may be more powerful than necessary for the generative discriminative framework, though further testing is necessary to confirm this.

6 Summary and conclusion

I've used this project to combine both regional and keypoint descriptor data to form a bag of words classifier.

My experimental results confirm that very disparate kinds of features can be effectively combined by the generative discriminative framework to improve performance. Indeed, probably the most effective use of the framework is to combine feature types that have as little overlap as possible in

what they describe.

7 Compiling and Running The Executable

See the README file included with the source for instructions on compiling and running the object recognition program.

References

- [1] Y. Li, L. G. Shaprio, and J. Bilmes *A Generative/Discriminative Learning Algorithm for Image Classification*. Department of Computer Science and Engineering, Department of Electrical Engineering, University of Washington
- [2] David G. Lowe *Distinctive Image Features from Scale-Invariant Keypoints* Computer Science Department University of British Columbia Vancouver, B.C., Canada January 5, 2004